



Managing disks

ONTAP 9

NetApp
December 01, 2021

Table of Contents

- Managing disks 1
 - When you need to update the Disk Qualification Package 1
 - How hot spare disks work 1
 - How low spare warnings can help you manage your spare disks 1
 - Display disk and partition ownership 2
 - Manually assign ownership of partitioned disks 3
 - Additional root-data partitioning management options 5
 - Configure automatic assignment of disk ownership 6
 - Which disk autoassignment policy to use 6
 - Remove a failed disk 7
 - Remove ownership from a disk 8
 - Sanitize a disk in maintenance mode 9
 - Set up an active-passive configuration on nodes using root-data partitioning 10
 - Set up an active-passive configuration on nodes using root-data-data partitioning 13
 - Commands for managing disks 15
 - Commands for displaying space usage information 17
 - Commands for displaying information about storage shelves 17

Managing disks

When you need to update the Disk Qualification Package

The Disk Qualification Package (DQP) adds full support for newly qualified drives. Before you update drive firmware or add new drive types or sizes to a cluster, you must update the DQP. A best practice is to also update the DQP regularly; for example, every quarter or semi-annually.

You need to download and install the DQP in the following situations:

- Whenever you add a new drive type or size to the node

For example, if you already have 1-TB drives and add 2-TB drives, you need to check for the latest DQP update.

- Whenever you update the disk firmware
- Whenever newer disk firmware or DQP files are available
- Whenever you upgrade to a new version of ONTAP.

The DQP is not updated as part of an ONTAP upgrade.

Related information

[NetApp Downloads: Disk Qualification Package](#)

[NetApp Downloads: Disk Drive Firmware](#)

How hot spare disks work

A hot spare disk is a disk that is assigned to a storage system and is ready for use, but is not in use by a RAID group and does not hold any data.

If a disk failure occurs within a RAID group, the hot spare disk is automatically assigned to the RAID group to replace the failed disks. The data of the failed disk is reconstructed on the hot spare replacement disk in the background from the RAID parity disk. The reconstruction activity is logged in the `/etc/message` file and an AutoSupport message is sent.

If the available hot spare disk is not the same size as the failed disk, a disk of the next larger size is chosen and then downsized to match the size of the disk that it is replacing.

How low spare warnings can help you manage your spare disks

By default, warnings are issued to the console and logs if you have fewer than one hot spare drive that matches the attributes of each drive in your storage system. You can change the threshold value for these warning messages to ensure that your system adheres to best practices.

You should set the `min_spare_count` RAID option to 2 to ensure that you always have the minimum recommended number of spare disks. You can use the `storage raid-options modify -node nodename -name option_name -value 2` to set the option.

Display disk and partition ownership

You can view disk ownership to determine which node controls the storage. You can also view the partition ownership on systems that use shared disks.

Steps

1. Display the ownership of physical disks using the `storage disk show -ownership` command:

```
cluster::> storage disk show -ownership
Disk      Aggregate Home      Owner    DR Home  Home ID      Owner ID      DR
Home ID   Reserver    Pool
-----
-----
1.0.0     aggr0_2    node2    node2    -        2014941509  2014941509    -
2014941509 Pool0
1.0.1     aggr0_2    node2    node2    -        2014941509  2014941509    -
2014941509 Pool0
1.0.2     aggr0_1    node1    node1    -        2014941219  2014941219    -
2014941219 Pool0
1.0.3     -          node1    node1    -        2014941219  2014941219    -
2014941219 Pool0
...
```

2. If you have a system that uses shared disks, display the partition ownership using the `storage disk show -partition-ownership` command:

```
cluster::> storage disk show -partition-ownership
```

Container		Root			Data		
Disk	Aggregate	Root Owner	Owner ID	Data Owner	Owner ID	Owner	
Owner ID							
1.0.0	-	node1	1886742616	node1	1886742616	node1	
1886742616							
1.0.1	-	node1	1886742616	node1	1886742616	node1	
1886742616							
1.0.2	-	node2	1886742657	node2	1886742657	node2	
1886742657							
1.0.3	-	node2	1886742657	node2	1886742657	node2	
1886742657							
...							

Manually assign ownership of partitioned disks

Manually assign ownership of partitioned disks overview

You can set the ownership of the container disk or the partitions manually or by using auto-assignment—just as you do for unpartitioned disks.



If a container disk fails in a half-populated shelf and is replaced, ONTAP will not auto-assign ownership. In this case, any assignment of new disks will need to be done manually. To make auto-assign work on half-populated shelves, place disks equally on lower half and 6 on far right bays to begin with. That is, 6 disks from bays 0-5 and 6 disks from bays 18-23. After the container disk is assigned in an ADP-configured system, ONTAP's software will handle any partitioning and partition assignments that are required, without user intervention.

Manually assign disks with root-data partitioning

For root-data partitioning there are three owned entities (the container disk and the two partitions) collectively owned by the HA pair.

About this task

The container disk and the two partitions do not all need to be owned by the same node in the HA pair as long as they are all owned by one of the nodes in the HA pair. However, when you use a partition in an aggregate, it must be owned by the same node that owns the aggregate.

Steps

1. Display the current ownership for the partitioned disk:

```
storage disk show -disk disk_name -partition-ownership
```

2. Set the CLI privilege level to advanced:

```
set -privilege advanced
```

3. Enter the appropriate command, depending on which ownership entity you want to assign ownership for:

If you want to assign ownership for the...	Use this command...
Container disk	<code>storage disk assign -disk disk_name -owner owner_name</code>
Data partition	<code>storage disk assign -disk disk_name -owner owner_name -data true</code>
Root partition	<code>storage disk assign -disk disk_name -owner owner_name -root true</code>

If any of the ownership entities are already owned, then you must include the `-force` option.

Manually assign disks with root-data-data partitioning

For root-data-data partitioning there are four owned entities (the container disk and the three partitions) collectively owned by the HA pair.

About this task

Root-data-data partitioning creates one small partition as the root partition and two larger, equally sized partitions for data.

Parameters must be used in the `disk assign` command to assign the proper partition of a root-data-data partitioned disk. You cannot use these parameters with disks that are part of a storage pool. The default value is false.

- The `[-data1 [true]]` parameter assigns the `data1` partition of a root-data1-data2 partitioned disk.
- The `[-data2 [true]]` parameter assigns the `data2` partition of a root-data1-data2 partitioned disk.

Steps

1. Display the current ownership for the partitioned disk:

```
storage disk show -disk disk_name -partition-ownership
```

2. Set the CLI privilege level to advanced:

```
set -privilege advanced
```

3. Enter the appropriate command, depending on which ownership entity you want to assign ownership for:

If you want to assign ownership for the...	Use this command...
Container disk	<code>storage disk assign -disk disk_name -owner owner_name</code>
Data1 partition	<code>storage disk assign -disk disk_name -owner owner_name-data1 true</code>
Data2 partition	<code>storage disk assign -disk disk_name -owner owner_name-data2 true</code>
Root partition	<code>storage disk assign -disk disk_name -owner owner_name -root true</code>

If any of the ownership entities are already owned, then you must include the `-force` option.

Additional root-data partitioning management options

Beginning with ONTAP 9.2, a new root-data partitioning option is available from the Boot Menu that provides additional management features for disks that are configured for root-data partitioning.

The following management features are available under the Boot Menu Option 9.

- Unpartition all disks and remove their ownership information

This option is useful if your system is configured for root-data partitioning and you need to reinitialize it with a different configuration.

- Clean configuration and initialize node with partitioned disks

This option is useful for the following:

- Your system is not configured for root-data partitioning and you would like to configure it for root-data partitioning
- Your system is incorrectly configured for root-data partitioning and you need to correct it
- You have an AFF platform or a FAS platform with only SSDs attached that is configured for the previous version of root-data partitioning and you want to upgrade it to the newer version of root-data partitioning to gain increased storage efficiency
- Clean configuration and initialize node with whole disks

This option is useful if you need to:

- Unpartition existing partitions
- Remove local disk ownership
- Reinitialize your system with whole disks using RAID-DP

Configure automatic assignment of disk ownership

You can configure ONTAP to automatically assign disk ownership according to a disk's stack, shelf, or bay. If configured, automatic disk ownership assignments occur 10 minutes after system initialization and every five minutes during normal system operation.

What you'll need

- Your system must adhere to the requirements for automatic disk ownership.
- If you have multiple stacks or shelves that must have different ownership, one disk must have been manually assigned on each stack or shelf so that automatic ownership assignment works on each stack or shelf.
- Use the `bay autoassign-policy` only for entry level platforms. If you try to use the `bay autoassign-policy` for a non-entry level platform, it will fail.

About this task

The behavior of the default automatic assignment policy depends on the system model. For entry level models, the default policy is equivalent to the `bay` policy. For all other systems, it is equivalent to the `stack` policy.

Steps

1. Configure automatic disk assignment:

```
storage disk option modify -autoassign-policy autoassign_policy -node  
node_name
```

- Use `stack` as the `autoassign_policy` to configure automatic ownership at the stack or loop level.
- Use `shelf` as the `autoassign_policy` to configure automatic ownership at the shelf level.
- Use `bay` as the `autoassign_policy` to configure automatic ownership at the bay level.

2. Verify the automatic assignment settings for the disks:

```
storage disk option show
```

```
cluster1::> storage disk option show
```

Node	BKg. FW. Upd.	Auto Copy	Auto Assign	Auto Assign Policy
-----	-----	-----	-----	-----
cluster1-1	on	on	on	default
cluster1-2	on	on	on	default

Which disk autoassignment policy to use

You can typically use the default autoassignment policy, which is equivalent to the `stack` policy for most systems, and to the `bay` policy for entry-level systems (AFF A2xx, FAS2xxx). However, for some configurations, you might need to change the

autoassignment policy.

You must select the appropriate autoassignment based on your configuration:

If you are using...	Then use the autoassignment policy value of...
Stand-alone entry-level system	stack
Entry-level systems in an HA configuration with a single, shared shelf	bay
Entry-level systems in an HA configuration with one stack of two or more shelves	shelf
MetroCluster configurations with one stack per node, two or more shelves	shelf
All other configurations	stack

Remove a failed disk

A disk that is completely failed is no longer counted by ONTAP as a usable disk, and you can immediately disconnect the disk from the disk shelf. However, you should leave a partially failed disk connected long enough for the Rapid RAID Recovery process to complete.

About this task

If you are removing a disk because it has failed or because it is producing excessive error messages, you should not use the disk again in this or any other storage system.

Steps

1. Find the disk ID of the failed disk:

```
storage disk show -broken
```

If the disk does not appear in the list of failed disks, it might be partially failed, with a Rapid RAID Recovery in process. In this case, you should wait until the disk is present in the list of failed disks (which means that the Rapid RAID Recovery process is complete) before removing the disk.

2. Determine the physical location of the disk you want to remove:

```
storage disk set-led -action on -disk disk_name 2
```

The fault LED on the face of the disk is lit.

3. Remove the disk from the disk shelf, following the instructions in the hardware guide for your disk shelf model.

Remove ownership from a disk

ONTAP writes disk ownership information to the disk. Before you remove a spare disk or its shelf from a node, you should remove its ownership information so that it can be properly integrated into another node.

What you'll need

The disk you want to remove ownership from must meet the following requirements:

- It must be a spare disk.

You cannot remove ownership from a disk that is being used in an aggregate.

- It cannot be in the maintenance center.
- It cannot be undergoing sanitization.
- It cannot be failed.

It is not necessary to remove ownership from a failed disk.

About this task

If you have automatic disk assignment enabled, ONTAP could automatically reassign ownership before you remove the disk from the node. For this reason, you disable automatic ownership assignment until the disk is removed, and then reenabling it.

Steps

1. If disk ownership automatic assignment is on, turn it off:

```
storage disk option modify -node node_name -autoassign off
```

2. If needed, repeat the previous step for the node's HA partner.
3. Remove the software ownership information from the disk:

```
storage disk removeowner disk_name
```

To remove ownership information from multiple disks, use a comma-separated list:

```
storage disk removeowner sys1:0a.23,sys1:0a.24,sys1:0a.25
```

4. If the disk is partitioned for root-data partitioning, remove ownership from the partitions by entering both of the following commands:

```
storage disk removeowner -disk disk_name -root true
```

```
storage disk removeowner -disk disk_name -data true
```

Both partitions are no longer owned by any node.

5. If you turned off disk ownership automatic assignment previously, turn it on after the disk has been removed or reassigned:

```
storage disk option modify -node node_name -autoassign on
```

6. If needed, repeat the previous step for the node's HA partner.

Sanitize a disk in maintenance mode

Beginning in ONTAP 9.6, you can perform disk sanitization in maintenance mode. This allows you to remove data from a disk or a set of disks on decommissioned or inoperable systems so that the data can never be recovered.

What you'll need

- The disk cannot be a self-encrypting disk (SED).

You must use the `storage encryption disk sanitize` command to sanitize an SED.

Encryption of data at rest

Steps

1. Boot into maintenance mode.
2. If the disks you want to sanitize are partitioned, unpartition each disk:

```
disk unpartition disk_name
```

3. Sanitize the specified disks:

```
disk sanitize start [-p pattern1|-r [-p pattern2|-r [-p pattern3|-r]]] [-c cycle_count] disk_list
```



Do not turn off power to the node, disrupt the storage connectivity, or remove target disks while sanitizing. If sanitizing is interrupted during the formatting phase, the formatting phase must be restarted and allowed to finish before the disks are sanitized and ready to be returned to the spare pool. If you need to abort the sanitization process, you can do so by using the `disk sanitize abort` command. If the specified disks are undergoing the formatting phase of sanitization, the abort does not occur until the phase is complete.

`-p pattern1 -p pattern2 -p pattern3` specifies a cycle of one to three user-defined hex byte overwrite patterns that can be applied in succession to the disks being sanitized. The default pattern is three passes, using 0x55 for the first pass, 0xaa for the second pass, and 0x3c for the third pass.

`-r` replaces a patterned overwrite with a random overwrite for any or all of the passes.

`-c cycle_count` specifies the number of times that the specified overwrite patterns are applied. The default value is one cycle. The maximum value is seven cycles.

`disk_list` specifies a space-separated list of the IDs of the spare disks to be sanitized.

4. If desired, check the status of the disk sanitization process:

```
disk sanitize status [disk_list]
```

5. After the sanitization process is complete, return the disks to spare status for each disk:

```
disk sanitize release disk_name
```

6. Exit maintenance mode.

Set up an active-passive configuration on nodes using root-data partitioning

When an HA pair is configured to use root-data partitioning by the factory, ownership of the data partitions is split between both nodes in the pair, for use in an active-active configuration. If you want to use the HA pair in an active-passive configuration, you must update partition ownership before creating your data aggregate.

What you'll need

- You should have decided which node will be the active node and which node will be the passive node.
- Storage failover must be configured on the HA pair.

About this task

This task is performed on two nodes: Node A and Node B.

All commands are input at the clustershell.

This procedure is designed for nodes for which no data aggregate has been created from the partitioned disks.

Steps

1. View the current ownership of the data partitions:

```
storage aggregate show-spare-disks
```

You can see that half of the data partitions are owned by one node and half are owned by the other node. All of the data partitions should be spare.

```
cluster1::> storage aggregate show-spare-disks

Original Owner: cluster1-01
Pool0
Partitioned Spares

Local
Local
Root Physical
Disk
Usable      Size
-----
1.0.0
0B 828.0GB
1.0.1
73.89GB 828.0GB
1.0.5
Type      RPM Checksum      Usable
BSAS      7200 block          753.8GB
BSAS      7200 block          753.8GB
BSAS      7200 block          753.8GB
```

```

0B 828.0GB
  1.0.6                BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.10               BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.11               BSAS    7200 block    753.8GB
0B 828.0GB

Original Owner: cluster1-02
Pool0
Partitioned Spares

Local
Local
Data
Root Physical
Disk
Usable      Size      Type      RPM Checksum    Usable
-----
-----
  1.0.2                BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.3                BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.4                BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.7                BSAS    7200 block    753.8GB
0B 828.0GB
  1.0.8                BSAS    7200 block    753.8GB
73.89GB 828.0GB
  1.0.9                BSAS    7200 block    753.8GB
0B 828.0GB
12 entries were displayed.

```

2. Enter the advanced privilege level:

```
set advanced
```

3. For each data partition owned by the node that will be the passive node, assign it to the active node:

```
storage disk assign -force -data true -owner active_node_name -disk disk_name
```

You do not need to include the partition as part of the disk name.

You would enter a command similar to the following example for each data partition you need to reassign:

```
storage disk assign -force -data true -owner cluster1-01 -disk 1.0.3
```

4. Confirm that all of the partitions are assigned to the active node.

```
cluster1::*> storage aggregate show-spare-disks
```

```
Original Owner: cluster1-01
```

```
Pool0
```

```
Partitioned Spares
```

```
Local
```

```
Local
```

```
Data
```

```
Root Physical
```

Disk	Type	RPM	Checksum	Usable
Usable	Size			

1.0.0	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.1	BSAS	7200	block	753.8GB
73.89GB 828.0GB				
1.0.2	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.3	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.4	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.5	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.6	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.7	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.8	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.9	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.10	BSAS	7200	block	753.8GB
0B 828.0GB				
1.0.11	BSAS	7200	block	753.8GB
0B 828.0GB				

```
Original Owner: cluster1-02
```

```
Pool0
```

```
Partitioned Spares
```

```
Local
```

```
Local
```

```
Data
```

```
Root Physical
```

Disk	Type	RPM	Checksum	Usable
Usable	Size			
-----	-----	-----	-----	-----
1.0.8	BSAS	7200	block	0B
73.89GB	828.0GB			

13 entries were displayed.

Note that cluster1-02 still owns a spare root partition.

- Return to administrative privilege:

```
set admin
```

- Create your data aggregate, leaving at least one data partition as spare:

```
storage aggregate create new_aggr_name -diskcount number_of_partitions -node
active_node_name
```

The data aggregate is created and is owned by the active node.

Set up an active-passive configuration on nodes using root-data-data partitioning

When an HA pair is configured to use root-data-data partitioning by the factory, ownership of the data partitions is split between both nodes in the pair, for use in an active-active configuration. If you want to use the HA pair in an active-passive configuration, you must update partition ownership before creating your data aggregate.

What you'll need

- You should have decided which node will be the active node and which node will be the passive node.
- Storage failover must be configured on the HA pair.

About this task

This task is performed on two nodes: Node A and Node B.

All commands are input at the clustershell.

This procedure is designed for nodes for which no data aggregate has been created from the partitioned disks.

Steps

- View the current ownership of the data partitions:

```
storage aggregate show-spare-disks -original-owner passive_node_name -fields
local-usable-data1-size, local-usable-data2-size
```

You should see that half of the data partitions are owned by one node and half are owned by the other node. All of the data partitions should be spare.

2. Enter the advanced privilege level:

```
set advanced
```

3. For each data1 partition owned by the node that will be the passive node, assign it to the active node:

```
storage disk assign -force -data1 -owner active_node_name -disk disk_name
```

You do not need to include the partition as part of the disk name

4. For each data2 partition owned by the node that will be the passive node, assign it to the active node:

```
storage disk assign -force -data2 -owner active_node_name -disk disk_name
```

You do not need to include the partition as part of the disk name

5. Confirm that all of the partitions are assigned to the active node:

```
storage aggregate show-spare-disks
```

```
cluster1::*> storage aggregate show-spare-disks
```

```
Original Owner: cluster1-01
```

```
Pool0
```

```
Partitioned Spares
```

```
Local
```

```
Local
```

```
Data
```

```
Root Physical
```

```
Disk
```

```
Type
```

```
RPM Checksum
```

```
Usable
```

```
Usable      Size
```

```
-----  
-----  
1.0.0          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.1          BSAS      7200 block      753.8GB  
73.89GB  828.0GB  
1.0.2          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.3          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.4          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.5          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.6          BSAS      7200 block      753.8GB  
0B  828.0GB  
1.0.7          BSAS      7200 block      753.8GB  
0B  828.0GB
```


If you want to...	Use this command...
Display a list of spare disks, including partitioned disks, by owner	<code>storage aggregate show-spare-disks</code>
Display the disk RAID type, current usage, and RAID group by aggregate	<code>storage aggregate show-status</code>
Display the RAID type, current usage, aggregate, and RAID group, including spares, for physical disks	<code>storage disk show -raid</code>
Display a list of failed disks	<code>storage disk show -broken</code>
Display the pre-cluster (nodescope) drive name for a disk	<code>storage disk show -primary-paths (advanced)</code>
Illuminate the LED for a particular disk or shelf	<code>storage disk set-led</code>
Display the checksum type for a specific disk	<code>storage disk show -fields checksum-compatibility</code>
Display the checksum type for all spare disks	<code>storage disk show -fields checksum-compatibility -container-type spare</code>
Display disk connectivity and placement information	<code>storage disk show -fields disk,primary-port,secondary-name,secondary-port,shelf,bay</code>
Display the pre-cluster disk names for specific disks	<code>storage disk show -disk diskname -fields diskpathnames</code>
Display the list of disks in the maintenance center	<code>storage disk show -maintenance</code>
Display SSD wear life	<code>storage disk show -ssd-wear</code>
Unpartition a shared disk	<code>storage disk unpartition (available at diagnostic level)</code>
Zero all non-zeroed disks	<code>storage disk zerospares</code>
Stop an ongoing sanitization process on one or more specified disks	<code>system node run -node nodename -command disk sanitize</code>
Display storage encryption disk information	<code>storage encryption disk show</code>

If you want to...	Use this command...
Retrieve authentication keys from all linked key management servers	<code>security key-manager restore</code>

Related information

[ONTAP 9 commands](#)

Commands for displaying space usage information

You use the `storage aggregate` and `volume` commands to see how space is being used in your aggregates and volumes and their Snapshot copies.

To display information about...	Use this command...
Aggregates, including details about used and available space percentages, Snapshot reserve size, and other space usage information	<code>storage aggregate show storage aggregate show-space -fields snap-size-total,used-including-snapshot-reserve</code>
How disks and RAID groups are used in an aggregate, and RAID status	<code>storage aggregate show-status</code>
The amount of disk space that would be reclaimed if you deleted a specific Snapshot copy	<code>volume snapshot compute-reclaimable</code>
The amount of space used by a volume	<code>volume show -fields size,used,available,percent-used volume show-space</code>
The amount of space used by a volume in the containing aggregate	<code>volume show-footprint</code>

Related information

[ONTAP 9 commands](#)

Commands for displaying information about storage shelves

You use the `storage shelf show` command to display configuration and error information for your disk shelves.

If you want to display...	Use this command...
General information about shelf configuration and hardware status	<code>storage shelf show</code>

If you want to display...	Use this command...
Detailed information for a specific shelf, including stack ID	<code>storage shelf show -shelf</code>
Unresolved, customer actionable, errors by shelf	<code>storage shelf show -errors</code>
Bay information	<code>storage shelf show -bay</code>
Connectivity information	<code>storage shelf show -connectivity</code>
Cooling information, including temperature sensors and cooling fans	<code>storage shelf show -cooling</code>
Information about I/O modules	<code>storage shelf show -module</code>
Port information	<code>storage shelf show -port</code>
Power information, including PSUs (power supply units), current sensors, and voltage sensors	<code>storage shelf show -power</code>

Related information

[ONTAP 9 commands](#)

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.